Progress Report 1

Implementation start-up (Week 1 of 6)

Cole Rodney Migisha	b29657	Charis Opol	b29654
Ainomugisha Faith	b27518	Lisa Lusinga Namugenyi	b30300
Ansasire Nehum Naamaye	b29785	Wakisa Isabelle	b30311

Project: LipScribe – Lip-Assisted Speech Recognition for Robust Live Captioning 20 October 2026

---

1 Summary

The project set-up is complete and implementation has started. We have created the Trello board that tracks the Work Breakdown Structure (WBS), created the project repository, downloaded the dataset and secured a GPU for training. As of Monday 19 October 2026 the team is implementing the system according to the WBS, starting with Week 1 (*Setup & Data*), which ends with Milestone M1 on Sunday 25 October.

2 Work completed

Trello board

The board turns the WBS into cards: one list per week, plus *In Progress*, *In Review* and *Done*. Each card has an owner, a reviewer, a due date and a checklist, and each week ends with a milestone card (M1–M9). Reference cards hold the team roles, the working agreement and the order in which work is dropped if the schedule slips.

Code repository

The GitHub repository (`https://github.com/coleRodneyMigisha/cv_project_LipScribe`) has been created for all code and documentation. All code is merged through pull requests that the card's reviewer must approve.

Dataset

The GRID audio-visual corpus has been downloaded. The data is stored outside Git and is not modified; integrity checks and the speaker-independent train/validation/test split follow this week.

GPU for training

A GPU has been secured for training the models, which removes the main risk identified in the implementation plan.

3 Week 1 plan and status

Week 1 tasks from the WBS, with their status on 20 October:

Card	Task	Owner	Due	Status
W1-01	Kick-off meeting, roles and Trello board	Charis Opol	Mon 19 Oct	**Done**
W1-02	GitHub repository	Ansasire Nehum Naamaye	Tue 20 Oct	**Done**
W1-04	Secure GPU access and storage	Charis Opol	Wed 21 Oct	**Done**
W1-05	Download the GRID corpus	Ainomugisha Faith	Wed 21 Oct	**Done**
W1-03	Pinned Python environment	Ansasire Nehum Naamaye	Wed 21 Oct	Planned
W1-06	MUSAN noise pools	Lisa Lusinga Namugenyi	Thu 22 Oct	Planned
W1-11	Module interface specification	Wakisa Isabelle	Thu 22 Oct	Planned
W1-07	Dataset integrity checker	Ainomugisha Faith	Fri 23 Oct	Planned
W1-09	Transcript normalisation and tokeniser	Lisa Lusinga Namugenyi	Fri 23 Oct	Planned
W1-12	Test framework and CI	Ansasire Nehum Naamaye	Fri 23 Oct	Planned
W1-08	Speaker-independent split	Ainomugisha Faith	Sat 24 Oct	Planned
W1-10	Mouth-cropping prototype	Cole Rodney Migisha	Sat 24 Oct	Planned
M1	Milestone review	All	Sun 25 Oct	Planned

4 Next steps

1. Build the pinned environment and confirm it on every member's machine and on the GPU machine.
2. Run the integrity checks on the downloaded data and fix the speaker-independent split (4 test speakers, 3 validation, the rest for training).
3. Build the tokeniser, the noise pools and the mouth-cropping prototype.
4. Hold the M1 review on Sunday 25 October and plan Week 2 (preprocessing; the audio-only baseline starts training).

5 Risks and issues

No blocking issues so far. With the GPU secured, the main open risk is software compatibility (MediaPipe supports only some Python versions), which the environment check this week addresses. The schedule and completion criteria for all milestones are in the accompanying *Project Milestones* document.